

## DBreeze Database Benchmarking.

(dbreeze version 01.003.20120509; doc. version 01.002.20120509)



*Professional, open-source, embedded, NoSql (Key/Value storage, horizontal scaling), transactional, ACID-compliant, multi-threaded, object database system for C# .NET 3.5> MONO.*

**Copyright © 2012 dbreeze.tiesky.com / Alex Solovyov / Ivars Sudmalis**

DBreeze is licensed under GPLv2. It's free for open-source projects, private and educational usage. Free for governmental needs of all countries and social non-profit organizations. For commercial closed-source projects appropriate license must be obtained.

Please, notify us about our software usage, so we can evaluate and visualize its efficiency.

### Testing Hardware

#### **Processor:**

Name: Intel Core I7 CPU 860 @ 2.8GHz  
Max Clock Speed: 2794 MHz  
Fron-Side Bus: 133 MHz  
L2 Cache: 8192 kB  
L3 Cach: 0 kB  
Number of Cores: 4  
Number of Logocal Processors: 8

#### **OS:**

Microsoft Windows 7 Ultimate SP1 64-bit

#### **HDD:**

WDC WD10EARS-00Y5B1 ATA Device  
Total size: 931.5 GB  
Total Cylinders: 121601  
Total Sectors: 1953520065  
Total Tracks: 31008255  
Bytes per Sector: 512  
Sectors per Track: 63  
Tracks per Cylinder: 255

## Benchmarking.

All tests source code you can find in the Solution, project VisualTester, class Benchmark.cs

### Data Inserting. One Thread

#### Test 1 (private void TEST\_1())

##### *Bulk Insert*

*Inserting of 1MLN data with identity key where identity seed = 1. Commit after all inserts.*

```
for (int i = 0; i < 10000000; i++)
{
    tran.Insert<int, byte[]>("t1", i, null);
}
tran.Commit();
```

*Quantity: 1 MLN*

*Insert took: 9.2 sec*

*Table file size: 18 MB*

#### Test 2 (private void TEST\_2())

##### *Bulk Insert*

*Inserting of 1MLN / 100K / 10MLN data with the growing key. Commit after all inserts.*

```
DateTime dt = new DateTime(1970, 1, 1);
```

```
for (int i = 0; i < 10000000; i++)
{
    tran.Insert<DateTime, byte[]>("t2", dt, null);
    dt = dt.AddSeconds(7);
}
```

```
tran.Commit();
```

*Quantity: 1 MLN*

*Insert took: 9.9 sec*

*Table file size: 22 MB*

*Quantity: 100 000 (100K)*

*Insert took: 997 ms*

*Table file size: 2.2 MB*

*Quantity: 10 MLN*

*Insert took: 98 sec  
Table file size: 220 MB*

### **Test 3 (private void TEST\_3())**

*Inserting of 1MLN / 100K data with the growing key. Commit after **every** inserts.*

This type of test is not about the speed, but about the final table file size. Repeats standard accumulating data behaviour of the program.

```
DateTime dt = new DateTime(1970, 1, 1);
```

```
for (int i = 0; i < 1000000; i++)  
{  
    tran.Insert<DateTime, byte[]>("t2", dt, null);  
    dt = dt.AddSeconds(7);  
    tran.Commit();  
}
```

*Quantity: 1 MLN  
Insert took: 212 sec  
Table file size: 29 MB*

*Quantity: 100 000 (100K)  
Insert took: 19 sec  
Table file size: 2.9 MB*

### **Test 3\_2.1 (private void TEST\_3\_2())**

*Updating of 1MLN / 100K data with the growing key. Commit after **every** update.  
(Initial data from TEST 3)*

```
for (int i = 0; i < 1000000; i++)  
{  
    tran.Insert<DateTime, byte[]>("t2", dt, null);  
    dt = dt.AddSeconds(7);  
    tran.Commit();  
}
```

*Quantity: 1 MLN  
Update took: 356 sec  
Table file size: 29 MB*

*Quantity: 100 000 (100K)  
Update took: 34 sec  
Table file size: 2.9 MB*

### **Test 3\_2.2 (private void TEST\_3\_2())**

*Bulk update of 1MLN / 100K.*

*Updating of 1MLN / 100K data with the growing key. Commit after **all** updates.  
(Initial data from TEST 3)*

```
for (int i = 0; i < 1000000; i++)
{
    tran.Insert<DateTime, byte[]>("t2", dt, null);
    dt = dt.AddSeconds(7);
}
tran.Commit();
```

*Quantity: 1 MLN*

*Update took: 28 sec*

*Table file size: 29 MB*

*Quantity: 100 000 (100K)*

*Update took: 2.8 sec*

*Table file size: 2.9 MB*

### **Test 7 (private void TEST\_7())**

*Bulk insert of Random Keys of 1MLN / 100K.*

*Insert of 1MLN / 100K data with the **random** key. Commit after **all inserts**.*

**Note**, this test is not only for speed calculation, but for checking Table File size.

Here we emulating filling the table with random 4 bytes keys.

For bulk inserts better to use preliminary sorted (in memory) data in ascending order, to make application more performant.

```
int vl = 0;
Random rnd = new Random();
for (int i = 0; i < 1000000; i++)
{
    vl = rnd.Next(1000000);

    tran.Insert<int, byte[]>("t5", vl, null);
}
```

```
tran.Commit();
```

*Quantity: 1MLN*

*Insert took: 330 sec*

*Table file size: 21 MB*

*Quantity: 100 000 (100K)*

*Insert took: 25 sec*

*Table file size: 2.9 MB*

### Test 7.1

*Insert of 200K data with the **random** key. Commit after every insert.*

*Quantity: 200 000 (200K)*

*Insert took: 67 sec*

*Table file size: 5.5 MB*

### Data Fetching

Fetch tests will be based on such insert:

```
DateTime dt = new DateTime(1970, 1, 1);

        for (int i = 0; i < 10000000; i++)
    {
        tran.Insert<DateTime, byte[]>("t2", dt, null);
        dt = dt.AddSeconds(7);
    }

    tran.Commit();
```

We got 10 MLN of 8 bytes keys with a “good jump”

### Test 8.1 (based private void TEST\_8())

*Iterating **forward** from 1 key, taking 1MLN / 100K / 10K / 1K*

```
int cnt = 0;
foreach (var row in tran.SelectForward<DateTime, byte[]>("t2").Take(1000000))
{
    cnt++;
}
```

*Select 1MLN took: 6 sec (6055 ms)*

*Select 100K took: 638 ms*

*Select 10K took: 81 ms*

*Select 1K took: 22 ms*

### Test 8.2 (based private void TEST\_8())

*Iterating **forward** from 1 key, taking 1MLN / 100K / 10K / 1K, also **taking value***

```

int cnt = 0;
byte[] val=null;
foreach (var row in tran.SelectForward<DateTime, byte[]>("t2").Take(1000000))
{
    val = row.Value;
    cnt++;
}

```

*Select 1 MLN took: 10 sec (10581 ms)*

*Select 100K took: 1165 ms*

*Select 10K took: 130 ms*

*Select 1K took: 29 ms*

### **Test 8.3 (based private void TEST\_8())**

*Iterating **backward** from 1 key, taking 1MLN / 100K / 10K / 1K*

```

int cnt = 0;
foreach (var row in tran.SelectForward<DateTime, byte[]>("t2").Take(1000000))
{
    cnt++;
}

```

*Select 1MLN took: 6200 ms*

*Select 100K took: 630 ms*

*Select 10K took: 81 ms*

*Select 1K took: 23 ms*

### **Test 8.4 (based private void TEST\_8())**

*Iterating **backward** from 1 key, taking 1MLN / 100K / 10K / 1K, also **taking value***

```

int cnt = 0;
byte[] val=null;
foreach (var row in tran.SelectForward<DateTime, byte[]>("t2").Take(1000000))
{
    val = row.Value;
    cnt++;
}

```

*Select 1 MLN took: 10501 ms*

*Select 100K took: 1070 ms*

*Select 10K took: 130 ms*

*Select 1K took: 29 ms*

### **Test 8.5 (based private void TEST\_8\_5())**

*Iterating **forward** from the key **N**, taking 100K*

*Iterating **backward** from the key **N**, taking 100K - takes the same time as forward*

10 MLN keys contain DateTime from 01.01.1970 - 21.03.1972.

We will make 3 test, key is starting from

~25% - 01.07.1970

~50% - 01.06.1971

~75% - 01.01.1972

```
int cnt = 0;
```

```
byte[] val = null;
```

```
//dtSearch we will change for every test
```

```
DateTime dtSearch = new DateTime(1970,7,1);
```

```
//dtSearch = new DateTime(1971, 6, 1);
```

```
//dtSearch = new DateTime(1972, 1, 1);
```

```
foreach (var row in tran.SelectForwardStartFrom<DateTime, byte[]>("t2",  
dtSearch,true).Take(100000))
```

```
{
```

```
    val = row.Value;
```

```
    cnt++;
```

```
}
```

*Select 100K ~25% took: 1152 ms*

*Select 100K ~50% took: 1102 ms*

*Select 100K ~75% took: 1098 ms*

### **Test 8.6 (based private void TEST\_8\_6())**

*Iterating forward from the key N to key M, with value acquiring*

In here every found key must be compared with To Key before returning

10 MLN keys contain DateTime from 01.01.1970 - 21.03.1972.

We will make 3 test, key is starting from

~25% - 01.07.1970

~50% - 01.06.1971

~75% - 01.01.1972

```
int cnt = 0;
```

```
byte[] val = null;
```

```
DateTime dtSearch = new DateTime(1970, 7, 1);
```

```
//dtSearch = new DateTime(1971, 6, 1);
```

```
//dtSearch = new DateTime(1972, 1, 1);
```

```
DateTime dtSearchStop = dtSearch.AddMonths(1);
```

```
foreach (var row in tran.SelectForwardFromTo<DateTime, byte[]>("t2", dtSearch, true,  
dtSearchStop,true))
```

```
{
```

```
    val = row.Value;
```

```
    cnt++;
```

```
}
```

*Select 100K ~25% took: 8542 ms Returned count: 382628*

Select 100K ~50% took: 6092 ms Returned count: 370286  
Select 100K ~75% took: 4430 ms Returned count: 382629

Select 100K ~25% took: 1116 ms .Take(100000)  
Select 100K ~50% took: 1116 ms .Take(100000)  
Select 100K ~75% took: 1096 ms .Take(100000)

### Test 8.7 (based private void TEST\_8\_7())

*Iterating backward from the key N to key M, with value acquiring*

In here every found key must be compared with To Key before returning

10 MLN keys contain DateTime from 01.01.1970 - 21.03.1972.

We will make 3 test, key is starting from

~25% - 01.07.1970

~50% - 01.06.1971

~75% - 01.01.1972

```
int cnt = 0;  
byte[] val = null;
```

```
DateTime dtSearch = new DateTime(1970, 7, 1);  
//dtSearch = new DateTime(1971, 6, 1);  
//dtSearch = new DateTime(1972, 1, 1);  
DateTime dtSearchStop = dtSearch.AddMonths(-1);
```

```
foreach (var row in tran.SelectBackwardFromTo<DateTime, byte[]>("t2", dtSearch, true,  
dtSearchStop,true))  
{  
    val = row.Value;  
    cnt++;  
}
```

Select 100K ~25% took: 5060 ms Returned count: 370286  
Select 100K ~50% took: 7500 ms Returned count: 382629  
Select 100K ~75% took: 9081 ms Returned count: 382628

Select 100K ~25% took: 1111 ms .Take(100000)  
Select 100K ~50% took: 1131 ms .Take(100000)  
Select 100K ~75% took: 1125 ms .Take(100000)

### Test 8.8 (based private void TEST\_8\_8())

*SkipForward from start 3MLN, then take 100000, with values acquiring (skipped keys don't acquire values)*

```
foreach (var row in tran.SelectForwardSkip<DateTime, byte[]>("t2", 3000000).Take(100000))
```



```

{
    val = row.Value;
    cnt++;
}

```

*SkipSelect 3MLN / 100K took: 3202 ms*

*SkipForward from start 6MLN, then take 100000, **with values acquiring***

*SkipSelect 6MLN / 100K took: 5605 ms*

*SkipForward from start 6MLN, then take 100000, **with values acquiring***

*SkipSelect 9MLN / 100K took: 7828 ms*

*SkipBackward from start 3MLN, then take 100000, **with values acquiring***

*SkipSelect 3MLN / 100K took: 3283 ms*

*SkipBackward from start 6MLN, then take 100000, **with values acquiring***

*SkipSelect 6MLN / 100K took: 5545 ms*

*SkipBackward from start 3MLN, then take 100000, **with values acquiring***

*SkipSelect 9MLN / 100K took: 7876 ms*

*SkipBackward from start 3MLN, then take 100000, **without value acquiring***

*SkipSelect 9MLN / 100K took: 7451 ms*

### **Test 8.9 (based private void TEST\_8\_9())**

*SkipForwardFromKey, skip 100K / 1MLN, then take 100000, **with values acquiring (skipped keys don't acquire values)***

```

int cnt = 0;
byte[] val = null;
DateTime dtSearch = new DateTime(1970, 7, 1);
//dtSearch = new DateTime(1971, 6, 1);
//dtSearch = new DateTime(1972, 1, 1);
foreach (var row in
    tran.SelectForwardSkipFrom<DateTime, byte[]>("t2", dtSearch,
100000).Take(100000))
{
    val = row.Value;
    cnt++;
}

```

*SkipForwardFromKey ~25% key skip 100K the take 100 K **with values acquiring***

*took: 1199 ms*

*SkipForwardFromKey ~25% key skip 1MLN the take 100 K with values acquiring  
took: 1825 ms*

*SkipForwardFromKey ~50% key skip 100K the take 100 K with values acquiring  
took: 1155 ms*

*SkipForwardFromKey ~50% key skip 1MLN the take 100 K with values acquiring  
took: 1809 ms*

*SkipForwardFromKey ~75% key skip 100K the take 100 K with values acquiring  
took: 1149 ms*

*SkipForwardFromKey ~75% key skip 1MLN the take 100 K with values acquiring  
took: 733 ms*

*SkipBackwardFromKey, skip 100K / 1MLN, then take 100000, with values acquiring (skipped  
keys don't acquire values)*

*SkipBackwardFromKey ~25% key skip 100K the take 100 K with values acquiring  
took: 1171 ms*

*SkipBackwardFromKey ~25% key skip 1MLN the take 100 K with values acquiring  
took: 1805 ms*

*SkipBackwardFromKey ~50% key skip 100K the take 100 K with values acquiring  
took: 1155 ms*

*SkipBackwardFromKey ~50% key skip 1MLN the take 100 K with values acquiring  
took: 1786 ms*

*SkipBackwardFromKey ~75% key skip 100K the take 100 K with values acquiring  
took: 1206 ms*

*SkipBackwardFromKey ~75% key skip 1MLN the take 100 K with values acquiring  
took: 1786 ms*

### **Test 9 (based private void TEST\_1\_9())**

*Selecting random keys 10K/ 100K / 1MLN, with and without values acquiring*

This test is built based on Test 1 (where we have ints from 1 to 1000000) all random search keys exist in this example

```
Random rnd = new Random();  
int key = 0;  
byte[] val = null;
```

```

for (int i = 0; i < 100000; i++)
{
    key = rnd.Next(9999999);
    var row = tran.Select<int, byte[]>("t1", key);

    if (row.Exists)
        val = row.Value; //or remarked if without values (just keys)
}

```

*Random Select 10K took: 1234 ms with value acquiring*  
*Random Select 10K took: 1199 ms without value acquiring*

*Random Select 100K took: 12177 ms with value acquiring*  
*Random Select 100K took: 11726 ms without value acquiring*

*Random Select 1MLN took: 122577 ms with value acquiring*  
*Random Select 1MLN took: 116120 ms without value acquiring*

### **Data Inserting. Multi-Threading**

In this test we will write data from many threads in different way.

#### **Test 10.1 (based private void TEST\_3\_1\_STARTER\_2())**

```

DateTime dt = new DateTime(1970, 1, 1);
for (int i = 0; i < 10000; i++)
{
    tran.Insert<DateTime, byte[]>(tableName, dt, null);
    dt = dt.AddSeconds(7);
    //tran.Commit();
}

tran.Commit();

```

*\* We have empty db and then start 100 threads, which have to write into their own table 1K (1000) records and **Commit of every table after every insert***

*The whole operation was finished in 7229 ms, we got 100 tables and 1000 records in every of it (13.8K records/second)*

*\* We have empty db and then start 100 threads, which have to write into their own table 10K*

(10000) records and **Commit of every table after every insert**

*The whole operation was finished in 87229 ms (87 sec), we got 100 tables and 10000 records in every of it (12K records/second)*

\* We have empty db and then start 100 threads, which have to write into their own table 1K (1000) records and **Commit of every table after all records are inserted in this table**

*The whole operation was finished in 506 ms, we got 100 tables and 1000 records in every of it (200K records/second)*

\* We have empty db and then start 100 threads, which have to write into their own table 10K (10000) records and **Commit of every table after all records are inserted in this table**

*The whole operation was finished in 4043 ms, we got 100 tables and 10000 records in every of it (247K records/second)*

\* We have empty db and then start 1 threads, which have to write into their own table 1MLN (1000000) records and **Commit of every table after all records are inserted in this table**

*The whole operation was finished in 9887 ms, we got 1 tables and 100000 records in it (102K records/second)*

\* We have empty db and then start 3 threads, which have to write into their own table 1MLN (1000000) records and **Commit of every table after all records are inserted in this table**

*The whole operation was finished in 10702 ms, we got 3 tables and 100000 records in every of it (300K records/second)*

\* We have empty db and then start 6 threads, which have to write into their own table 1MLN (1000000) records and **Commit of every table after all records are inserted in this table**

*The whole operation was finished in 18290 ms, we got 3 tables and 100000 records in every of it (333K records/second)*

*After this test we have 6 files 22MB each, in total 132 MB. Writing speed was 7.3 MB/s and it was random writing speed.*

## **Multiple tables.**

This test is dedicated to the efficiency of working with multiple tables. For now every table

locates in 3 OS files. One for pure data, second for rollback data and third for efficient rollback usage. And this test was done under NTFS (which uses b+tree for locating files and files fragments).

We will try in following tests to create, to write and to read data from 34K/ 200K/1 MLN tables.

### Multiple tables. Write/Read.

#### Test 11 (based private void TEST\_9() - write and TEST\_9\_3() - read)

Database is empty to the beginning of the test. We will use such script to create and to write a small piece of data into 34K/ 200K/1 MLN tables:

*//This value will change from test to test*

*int totalTablesToCreate = 1000000;*

```
for (int i = 1; i <= totalTablesToCreate; i++)
{
    using (var tran = engine.GetTransaction())
    {
        try
        {
            tran.Insert<byte[], byte[]>("t" + i, new byte[] { 0 }, new byte[] { 0 });
            tran.Commit();
        }
        catch (Exception ex)
        {
            Console.WriteLine(ex.ToString());
        }
    }
}
```

and then we used such script to read data from all these tables:

*//This value will change from test to test*

*int totalTablesToCreate = 1000000;*

*byte[] v = null;*

*for (int i = 1; i <= totalTablesToCreate; i++)*

```
{
    using (var tran = engine.GetTransaction())
    {
        try
        {
```

```

        var row = tran.Select<byte[], byte[]>("t" + i, new byte[] { 0 });
        v = row.Value;
    }
    catch (Exception ex)
    {
        Console.WriteLine(ex.ToString());
    }
}
}

```

## Results:

### **Write + access time:**

34K tables - 33 sec ~ 1030 tables per second  
 200K tables - 263 sec ~ 760 tables per second  
 1MLN tables - 1360 sec ~ 735 tables per second

### *Filesystem information for 1MLN tables:*

*Quantity of files in DBreeze instance folder - 3.000.006 (together with scheme and transaction journal)*

*Total folder size: 140 MB*

*Total size which folder resides on the physical disk (depends upon cluster size): 7,67 GB*

### **Read + access time:**

34K tables - 15 sec ~ 2266 tables per second  
 200K tables - 121 sec ~ 1652 tables per second  
 1MLN tables - 736 sec ~ 1358 tables per second

### **Memory:**

*Write test, project started from 28 MB and finished with 33 MB in RAM.*

*Read test, project started from 28 MB and finished with 28 MB in RAM.*